

ΥΣ13 - Computer Security

Network Security

Κώστας Χατζηκοκολάκης

Context

- Computers connected in a network
 - but also: smartphones, fridges, IoT devices, ...
- Each device has an IP address
- Packets are routed via intermediate nodes
- Still using IPv4 (almost 40 year old!)
 - very very hard to replace

- **Attacker model**
 - Intercept packets
 - Modify packets
 - Inject packets
 - Participate in any protocol
- Useful to consider combinations of the above

Internet protocol suite

Four layers (7 in the OSI model)

- **Link**
 - Physical addresses
 - Physical aspects of communication
- **Internet**
 - Addressing (source/dest IP)
 - Routing
 - Time to live

Internet protocol suite

Four layers (7 in the OSI model)

- **Transport**

- Source/dest ports
- Ordering of packets (Sequence numbers)
- ACKs, checksums

- **Application**

- The “real data”, application-dependent

Internet protocol suite

Protocols

- **Link**

- Ethernet
- WiFi
- DSL
- ...

- **Internet**

- IP
- ICMP
- ...

Internet protocol suite

Protocols

- **Transport**

- TCP
- UDP
- ...

- **Application**

- HTTP / HTTPS
- SSH
- SMTP
- ...

Internet protocol suite

Packet example

Physical Layer: eth
the 2 MAC addresses
+ IP indication

Network Layer: IP
IP addresses, TTL,
checksum, fragmentation

Transport Layer: TCP
Ports, Seq Ack numbers,
checksum, timestamps

Application Layer: HTTP
Request: GET
Request URI
Referrer
User-agent info
Connection info

```
0000  00 0f db 4d 77 95 00 0d 93 b0 a3 24 08 00 45 00  ...Mw.....$.E.
0010  01 75 c8 de 40 00 40 06 44 dd c0 a8 01 2e 40 ec  ..u..@.D....@.
0020  29 05 e6 10 00 50 15 86 10 4d 25 b6 67 ed 80 18  )....P..M%.g...
0030  ff ff 2d 2f 00 00 01 01 08 0a 2f 41 cf 64 62 38  ..-./...../A.db8
0040  81 9a 47 45 54 20 2f 63 6e 6e 2f 32 30 30 36 2f  ..GET /cnn/2006/
0050  55 53 2f 30 32 2f 32 37 2f 6b 61 74 72 69 6e 61  US/02/27/katrina
0060  2e 70 6f 6c 6c 2f 74 31 2e 32 31 33 35 2e 6d 6f  .poll/tl.2135.mo
0070  6e 2e 62 65 61 64 73 2e 61 70 2e 6a 70 67 20 48  n.beads.ap.jpg H
0080  54 54 50 2f 31 2e 31 0d 0a 41 63 63 65 70 74 3a  TTP/1.1..Accept:
0090  20 2a 2f 2a 0d 0a 41 63 63 65 70 74 2d 4c 61 6e  /*..Accept-Lan
00a0  67 75 61 67 65 3a 20 65 6e 0d 0a 41 63 63 65 70  guage: en..Accep
00b0  74 2d 45 6e 63 6f 64 69 6e 67 3a 20 67 7a 69 70  t-Encoding: gzip
00c0  2c 20 64 65 66 6c 61 74 65 0d 0a 52 65 66 65 72  , deflate..Refer
00d0  65 72 3a 20 68 74 74 70 3a 2f 2f 77 77 77 2e 63  er: http://www.c
00e0  6e 6e 2e 63 6f 6d 2f 0d 0a 55 73 65 72 2d 41 67  nn.com/..User-Ag
00f0  65 6e 74 3a 20 4d 6f 7a 69 6c 6c 61 2f 35 2e 30  ent: Mozilla/5.0
0100  20 28 4d 61 63 69 6e 7a 6f 73 68 3b 20 55 3b 20  (Macintosh; U;
0110  50 50 43 20 4d 61 63 20 4f 53 20 58 3b 20 65 6e  PPC Mac OS X; en
0120  29 20 41 70 70 6c 65 57 65 62 4b 69 74 2f 34 31  ) AppleWebKit/41
0130  37 2e 39 20 28 4b 48 54 4d 4c 2c 20 6c 69 6b 65  7.9 (KHTML, like
0140  20 47 65 63 6b 6f 29 20 53 61 66 61 72 69 2f 34  Gecko) Safari/4
0150  31 37 2e 38 0d 0a 43 6f 6e 6e 65 63 74 69 6f 6e  17.8..Connection
0160  3a 20 6b 65 65 70 2d 61 6c 69 76 65 0d 0a 48 6f  : keep-alive..Ho
0170  73 74 3a 20 69 2e 61 2e 63 6e 6e 2e 6e 65 74 0d  st: i.a.cnn.net.
0180  0a 0d 0a  ...
```


- Connectionless communication
 - using only source/dest IP addresses
- Routing
 - communication across network boundaries
 - routing tables kept by routers
 - no authentication
- Fragmentation & reassembly
 - No reliability

- Connection-based communication
 - identified by source/dest IP + port (multiplexing)
- Server process “listens” to a port
 - Often determined by the application protocol (HTTP, SMTP, etc)
- Client process connects to dest IP+port
 - Source port selection usually random
- Connection established by handshake
- Reliability

UDP

- Connectionless communication over IP
- Fast alternative to TCP
 - Only 8 bytes overhead, no handshakes
 - Stateless
- Some higher-level features
 - addressing based on IP+port (multiplexing)
 - checksums
- But many missing
 - No ACKs (unreliable)
 - No ordering
- Often used for “streaming”-like applications

Traceroute

traceroute to google.com (216.58.215.46), 30 hops max, 60 byte packets


```
1  _gateway (195.134.67.1)  0.715 ms  0.789 ms  0.884 ms
2  uoa-ilisia-1-gw.kolettir.access-link.grnet.gr (62.217.96.172)  0.763 ms  0.796 ms  0.
3  grnet-ias-geant-gw.mx1.ath2.gr.geant.net (83.97.88.65)  1.574 ms  1.630 ms  1.620 ms
4  ae0.mx2.ath.gr.geant.net (62.40.98.140)  31.556 ms  31.650 ms  31.547 ms
5  ae2.mx1.mil2.it.geant.net (62.40.98.150)  25.654 ms  27.861 ms  27.793 ms
6  72.14.203.32 (72.14.203.32)  25.593 ms  25.766 ms  25.500 ms
7  108.170.245.73 (108.170.245.73)  64.548 ms 108.170.245.89 (108.170.245.89)  73.238 m
8  209.85.142.221 (209.85.142.221)  72.001 ms 72.14.238.21 (72.14.238.21)  71.999 ms 6
9  216.239.35.201 (216.239.35.201)  78.302 ms  78.299 ms  78.277 ms
10 209.85.251.217 (209.85.251.217)  54.466 ms 72.14.238.54 (72.14.238.54)  54.472 ms 1
11 108.170.245.1 (108.170.245.1)  52.509 ms  52.443 ms  50.669 ms
12 108.170.235.15 (108.170.235.15)  54.116 ms  51.975 ms  51.967 ms
13 par21s17-in-f14.1e100.net (216.58.215.46)  51.943 ms  54.241 ms  54.202 ms
```

Traceroute

- Time to live (TTL)
 - IP header
 - Decreased at every hop
 - If 0 the router discards and notifies the originator
- Traceroute: repeatedly send packets
 - with TTL = 1, 2, ...
 - 3 packets for every value
 - Until we reach the host (or a threshold)
 - Routers might not respond

TCP 3-way handshake


- Connection identified by **source/dest address/port**
- **Sequence numbers (SN)** in every message
- Handshake
 - SYN(SN_c)
 - SYN(SN_s)-ACK(SN_c)
 - ACK(SN_s)
 - Data-exchange (bidirect.)



Three-Way Handshake

TCP 3-way handshake


- Connection identified by **source/dest address/port**
- **Sequence numbers (SN)** in every message
- Handshake
 - SYN(SN_c)
 - SYN(SN_s)-ACK(SN_c)
 - ACK(SN_s)
 - Data-exchange (bidirect.)
- **What can go wrong here?**



Three-Way Handshake

SYN flood


- Flood the server with SYNs
- But **no ACK!**
- Connections stay “half-open” on the server until they timeout
 - Keeping **state** consumes **resources**
 - Can lead to **Denial of Service (DoS)**



TCP-SYN Flood Attack

SYN flood


- Flood the server with SYNs
- But **no ACK!**
- Connections stay “half-open” on the server until they timeout
 - Keeping **state** consumes **resources**
 - Can lead to **Denial of Service (DoS)**
- Can the server limit the number of SYNs from the **same host**?
 - **No!** the attacker can easily “**spoof**” the sender IP



TCP-SYN Flood Attack

IP spoofing


- Can we impersonate a client?



Three-Way Handshake

IP spoofing


- Can we impersonate a client?
 - Trivial if we control an intermediate router!
 - If we don't?



Three-Way Handshake

IP spoofing


- Can we impersonate a client?
 - Trivial if we control an intermediate router!
 - If we don't?
- We can still send packets with a spoofed IP, without access to the replies



Three-Way Handshake

IP spoofing

- Can we impersonate a client?
 - Trivial if we control an intermediate router!
 - If we don't?
- We can still send packets with a spoofed IP, without access to the replies
 - It's sufficient to guess SNs for the ACK!
 - A(C) → S : SYN(SN_a)
 - S → C : SYN(SN_s)-ACK(SN_a)
 - A(C) → S : ACK(SN_s)



Three-Way Handshake

IP spoofing

- Can we guess the server's SN?
- Initial Sequence Number
 - Counter incremented over time and for every new connection
 - Predictable!
- Routers expect ISN to be increasing
 - Protocol bugs are hard to fix (compared to implementation bugs)

IP spoofing

- Can we guess the server's SN?
- Initial Sequence Number
 - Counter incremented over time and for every new connection
 - Predictable!
- Routers expect ISN to be increasing
 - Protocol bugs are hard to fix (compared to implementation bugs)
- Solution?

IP spoofing

- Can we guess the server's SN?
- Initial Sequence Number
 - Counter incremented over time and for every new connection
 - Predictable!
- Routers expect ISN to be increasing
 - Protocol bugs are hard to fix (compared to implementation bugs)
- Solution?
 - Different ISN for each client! How?

IP spoofing

- Can we guess the server's SN?
- Initial Sequence Number
 - Counter incremented over time and for every new connection
 - Predictable!
- Routers expect ISN to be increasing
 - Protocol bugs are hard to fix (compared to implementation bugs)
- Solution?
 - Different ISN for each client! How?
 - RFC 6528:
ISN = Timer + PRF(localip, localport, remoteip, remoteport, secretkey)
 - Why we include secretkey?

IP spoofing

Why is it bad?

- Bypass IP-based authorization
 - Still **widely-used** today
 - SMTP, web-services, firewall IP white/black-listing, etc
- Inject data to existing connection
 - **DNS response** (UDP, no SN at all!)
- Reset existing connections (RST)
 - SNc is needed, but only approximately
 - Denial of service, or exploit to break some other protocol

Conclusion : For serious security we need to build **on top of TCP**

Denial of Service

- Remotely consume a resource of the server
 - Bandwidth,
 - CPU
 - Memory
 - ...
- Until the resource is depleted
 - no more clients can connect

Denial of Service

- Typically involves some sort of **flooding**
- SYN flooding

Denial of Service

- Typically involves some sort of **flooding**
- SYN flooding
- Ping flooding

- ICMP echo request

```
~$ ping google.com
```

```
PING google.com (216.58.215.46) 56(84) bytes of data.
```

```
64 bytes from par21s17-in-f14.1e100.net (216.58.215.46): icmp_seq=1 ttl=47 time=52
```

- Low-level protocol (no use of TCP), can send packets fast
- The server sends a reply back

Denial of Service

- Typically involves some sort of **flooding**

- SYN flooding

- Ping flooding

- ICMP echo request

```
~$ ping google.com
```

```
PING google.com (216.58.215.46) 56(84) bytes of data.
```

```
64 bytes from par21s17-in-f14.1e100.net (216.58.215.46): icmp_seq=1 ttl=47 time=52
```


- Low-level protocol (no use of TCP), can send packets fast
 - The server sends a reply back

- We need **more resources than the server**

- Use many senders at once


Smurf attack

- Send an Echo to a **broadcast** address
- **Spoof** the sender IP with the sender's
- All machines flood the victim



Distributed Denial of Service (DDoS)

- Compromise hosts via virus, worm, etc
- Coordinate the attack
- Hard to distinguish from legitimate users



Fork Bomb

- Another kind of Dos
- Fork, and keep forking in the children
 - exponential growth
- Consumes OS resources for process management
- Try this in your own machine!

```
~$ :(){ :|:& };:
```

```
# some other terminal
```


```
~$ ls
```

```
bash: fork: retry: Resource temporarily unavailable
```

Preventing SYN floods

Crucial properties

- Packets need to arrive from multiple source IPs
 - otherwise trivial to filter
- The adversary spoofs the sender IP
 - but does not get replies!
- The server needs to **keep state** for all fake clients



Three-Way Handshake


Preventing SYN floods

Crucial properties

- Packets need to arrive from multiple source IPs
 - otherwise trivial to filter
- The adversary spoofs the sender IP
 - but does not get replies!
- The server needs to **keep state** for all fake clients

Idea


- Make the client store the state!
- Only store state in the server for clients that have proven to get our replies



Three-Way Handshake

Preventing SYN floods

SYN cookies




Three-Way Handshake

Preventing SYN floods

SYN cookies

- Encode the state in the SNs send to client
 - Then **forget** about the connection (no state!)
- Check the SNs contained in the client's ACK
 - Store state only if ok
- Spoofing the source is useless
 - The adversary needs to control the users or the network




Three-Way Handshake

Preventing SYN floods

SYN cookies

- Encode the state in the SNs send to client
 - Then **forget** about the connection (no state!)
- Check the SNs contained in the client's ACK
 - Store state only if ok
- Spoofing the source is useless
 - The adversary needs to control the users or the network
- One approach
 - $SNs = H(\text{ports, ips, key, time}) \parallel \text{time}$




Three-Way Handshake

Preventing SYN floods

SYN cookies

- Encode the state in the SNs send to client
 - Then **forget** about the connection (no state!)
- Check the SNs contained in the client's ACK
 - Store state only if ok
- Spoofing the source is useless
 - The adversary needs to control the users or the network
- One approach
 - $SNs = H(\text{ports, ips, key, time}) \parallel \text{time}$
- Protocol compliant, but problematic if ACK is lost



Preventing DoS

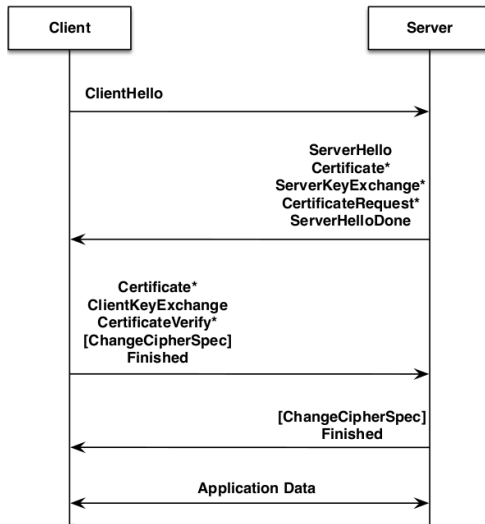
- Client needs to **solve a puzzle** to connect
 - Eg: brute-force a hash (within a controlled range of values)
- Generic solution, also used to prevent spam
- But requires changes to both client and server

Achieving secure communication

- TCP is an **inherently insecure** protocol
 - no security against an adversary who controls the network
 - limited security against an adversary who simply participates
- Solution
 - Use crypto to **build a secure connection** over an insecure network
- Most widely used: TLS
 - Also: IPSec, SSH, ...
- We can also tunnel the traffic of an **entire network**
 - Secure VPN

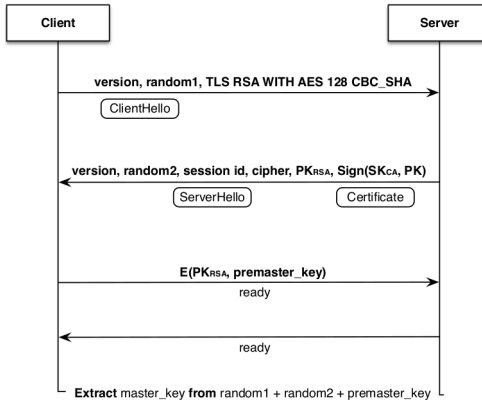
- Widely used in web-browsers
- Crucial use of crypto:
 - Assymmetric-crypt: exchange keys
 - Symmetric crypto: encrypt the main traffic
 - Digital signatures: authentication

TLS handshake



* = προαιρετικά

TLS handshake



References

- Ross Anderson, Security Engineering, Chapter 21
- A look back at "security problems in the TCP/IP protocol suite"
- SYN cookies
- Bypassing domain control verification with DNS response spoofing